Title of the article

Mahieddine Yaker
and Julien Cartigny
and Gilles Grimaud
IRCICA
University of Lille
address in Lille
Email: yyy@yyy.com

Chrystel Gaber
and Xiao Han
and Vicente Sanchez-Leighton
Orange Labs,
Châtillon, France
Email: firstname.lastname@orange.com

Abstract—Abstract

TBD

I. Introduction

TBD

mds
March 05, 2018

II. Internet of Things security issues

Oravec and co [1] described in their paper the increase of utilization of devices and applications dedicated to Internet of Things, particularly in the case of household appliances and house management. They wished to demonstrated the new issues that should be considered with the expansion of IoT into the house. One of their main concern is the security and privacy integrity of users. They descried scenario about different attacks or other information due to fallible household devices and non-secure softwares.

In that kind of environment, Sivaraman and co [2] described a scenario based on an attack on all house IoT devices via a smartphone application. By using the smartphone and network accesses into the house, they found all available IoT devices. In their paper, they show how it is possible to get control of each device. With this attack, they show that even if there is strong network security mechanisms (like router equipped with firewalls) if any component of the House IoT has an application level issue, the security and privacy is no more guaranteed into the house. This attack works with a smartphone, but we can also imagine a scenario where an another connected device (Webcam, connected TV..) can be an attack vector.

III. Isolation model

TBD by Lille

A. Memory isolation proto-kernel

PIP [3] is a protokernel: it allows for kernels, ranging from hypervisors to monolithic kernels, to be developed as user mode applications (figure 1). This means that only PIP is executed in kernel mode (i.e., the privileged mode of the hardware)(figure . Indeed, code running in kernel mode has direct access to the whole memory and hardware. It is thus


Fig. 1. Pip architecture view

clearly better, from a security point of view, to keep this code as minimal as possible. This stems from the general principle that the trusted computing base (TCB) should be kept minimal. One of the PIP key design is to have a minimal number of functionalities while ensuring strong isolation. More precisely, PIP only manages memory isolation and redirection of interruptions to user space code, and has only 10 system calls. Contrary to micro-kernel, it means that components like scheduler, IPC and authorization are not included in the PIP kernel, neither other mechanism available in monolithic kernel like device abstraction layer, file systems, network stack, etc.

PIP proposes a hierarchic memory isolation model (figure 2) to partition the available memory among several memory-limited space called partition. A partition is a set of physical memory pages mapped to virtual address. At start, all the available memory (as defined inside PIP configuration) makes up the ROOT partition. The root partition can include mapped registers from hardware devices, thus providing access to hardware functionalities. Code executed in a parent partition can request (via system calls to PIP) segregation to its memory space: a new child partition is created on which the parent partition can mapped a subset of its own mapped physical memory pages, thus *delegating* part of its memory space to the child partition. This model is recursive: any partition can create a child partition and delegate part of its memory space. While a partition can read and write in the memory of its
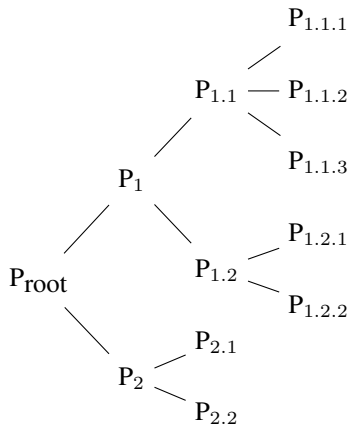
Fig. 2. The view by Pip of the partition tree

child partition (we call this property vertical sharing), sibling partitions cannot access each other's memory (we call this property horizontal isolation).
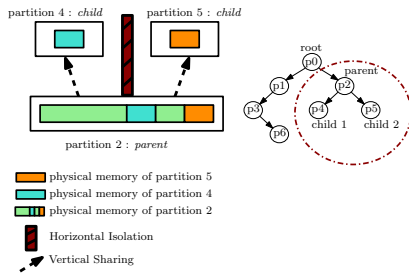
B. Pip memory management

Fig. 3. Isolation model

Pip model ensure formal isolation between partitions. This isolation is guaranteed by two properties :

- **Horizontal isolation** : Two distinct partitions (ex: *child 1* and *child 2* in figure 3) can not share the same memory pages.
- **Vertical sharing** : All pages used by a child are mapped into its parent.

As seen in figure 3, for partition creation, a parent gives the required amount of memory pages from its own pages to its child, and theses pages belong to this new partition.

In order to create partition and manage each child memory, Pip provide a set a function available for each partition.

- **createPartition** : create a new partition.
- **deletePartition** : delete a new partition.
- **addVAddr** : gives a page to a child partition (no more available into the parent).
- **removeVAddr** : reclaims a page from a child partition (no more available into the parent).

C. Interrupt management

To manage interruptions, PIP has two different behaviours depending of the kind of interruption:

- Software interrupts (fault or system calls) are relayed to the parent partition (with the exception of the PIP system calls).
- Hardware interrupts are relayed to the ROOT partition which is in charge to forward it to, for instance, a network card driver.

Furthermore, a parent receives all software interruptions from its children, thus controlling which interactions have its children with the system.

In addition, Pip API provides two function for software interrupt between partitions :

- **dispatch** : send an interrupt signal to a given partition.
- **resume** : restores a previous interrupted partition.

Any partition can send and receive a signal only from its parent, one of its child and from Pip. For example, in figure 3 There is no available signal transmission between *p1* and *p2* or *p3* and *p2*.

D. IPC mechanism

Fig. 4. Simple IPC mechanism

Pip does not provide IPC between two distinct partitions. However, this mechanism can be implemented in userland, by using software interrupt or Pip API dispatch function. For example (figure 4), if *Child 1* wants to send a message to *Child 2*, it has to send a dispatch to its *parent* (step A). The parent handles this software interrupt and transmit the signal to *Child 2* (step B).

For more complex architecture (like figure 2), this mechanism can be recursively used to send signal between any partitions (ex : P1.1 to P.1.2.2).

IV. INDUSTRIAL ECOSYSTEM

This section describes the industrial ecosystem and the needs for a security architecture based on isolation. First, we identify and define the stakeholders involved in the use and management of an IOT or M2M device. The responsibilities of each actor are then described and finally, the requirements for an isolation-based architecture are listed.

We consider that a domain is an isolated area which belongs to an entity. The notion of domain is precised in section V-A.

We distinguish device ressources and domain ressources. Device ressources correspond to the isolated partitions (domains) and the ressources exposed by software, actuators

or sensors drivers installed by the packager in the owner domain and which are related to the device. Domain ressources correspond to any ressource not related to the device and created within the given domain. En example of device ressource is the firmware or its updater. An example of a domain ressource is a key to create an SSL channel between a remote platform and a given domain.

A. Actors

We identify 7 stakeholders, namely the manufacturer, packager, maintainer, owner, administrator, service provider and user, which interact with the IOT ot M2M device. One device may be associated to 1 packager, 1 owner, 1 administrator and several manufacturers, maintainers and service providers. One actor may hold several roles. For example, one actor can be at the same time a manufacturer, a packager and a maintainer.

1) Manufacturer: Manufacturers provide at least one component of the device, either a physical component with its driver, an operating system or an isolation solution. Manufacturers also provide updates & patches for the component and the driver to the packager and/or maintainers.

2) Packager: It issues the device and the isolation solution. It also provides each device an identifier that will subsequently allow it to be identified as well as intial secrets that will allow the owner to access the device. It may also provide the drivers and software components for the sensors or actuators on the device.

3) Maintainer: It maintains the device, monitors the hardware components status and deploys the patches delivered by the manufacturer.

4) Owner: The device belongs to the Owner who defines an access policy to authorize which other entities are authorized to use or manage the device or domains within the device. The Owner can possess several devices. It receives initial credentials from the manufacturer that allows it to access each device in his fleet.

5) Administrator: It enforces the policy defined by the owner for his fleet of devices. If any modification, such as adding a new entity or modifying its granted permissions, is required, it requests a decision from the owner. This role can either be assumed or delegated to a third party by the owner.

6) Service Provider: It delivers a user-friendly service on one or multiple devices using the ressources authorized by the Owner. It installs or activates a service within the device. Muliple Service Providers can co-exist on the same device.

7) User: Users consume the services provided by the entities mentioned above. We can distinguish platform users and technical users. Service users subscribe to the services provided by the Service Providers. Technical users are typically members of the Owner, Manufacturer, Maintainer entities and they will perform authorized administration actions such as creating a new domain, granting rights to a new user, loading a service on the device. In the rest of this article, we focus on platform users.

B. Responsibilities model

The responsibilities described in this section are summarized in table I.

The packager gathers hardware, drivers and OS from various manufacturers and creates a device with this collection of components. It provides the mechanisms and temporary credentials to personalize the device. At delivery, it provides the keys of the owner domain. The packager is responsible for providing drivers which are compatible with virtualization and usage by multiple entities. For example, to control the position of an armed robot, it is better to literally express the coordinates of the destination. The command "go to (x=10;y=40;z=20)" leads to less confusion than "go to current position + (x=10; y=20; z=30). The packager provides a platform in a safe state to the owner and does not keep any access to the owner domain. In particular, the packager does not manage access control to the drivers.

The owner finalizes the personalization of the device after delivery by the manufacturer. In particular, the owner should modify its temporary credentials. Owner authorizes the usage of the device resources. If some resources are very sensitive, he provides the access to them through the use of a token which he delivers and verifies. The owner can choose that some resources are accessible freely without the use of a token and thus reducing the security. The owner must not have any control or visibility on the actions performed by other entities unless it touches sensitive functions which the owner has decided to control through token verification.

If the owner delegates his tasks to an administrator, then the administrator can configure which resources need to be accessed with a token and the administrator is responsible for verifying the token and should not be able to control or visualize the actions performed by other actors unless they concern his perimeter of action.

The maintainers keep the firmware, driver or any software in their perimeter up to date. Their role is to keep track of the latest updates and changes provided by manufacturers and deploy them on the field. Both responsibilities (correcting software and deploying updates) are key in the future as the regulators start to take actions on this point. For instance, the European Commission's overall security strategy [4] requires vendors to commit to update their software in the event of newly disclosed vulnerabilities, as part of "duty of care" principle. Such initiative also exists in the United States with the Internet Of Things (IoT) Cybersecurity Improvement Act of 2017 [5].

Any entity which owns a partition (administrator, service provider, and maintainer) has to authenticate the machine (or user) who is sending the commands to the partition. This entity can choose to not perform and access control verification, thus

TABLE I
RESPONSIBILITIES MATRIX

Responsibilities	Manufacturer	Packager	Maintainer	Owner	Administrator	Service Provider
Provide device components (hardware & drviers, OS, isolation solution)	X					
Gather components into a device		X				
Use or provide virtualization & multi-tenant compatible components		X				
Provide mechanisms & temporary credentials for initial device personalization before delivery to the owner		X				
Provide temporary credentials for initial domain personalization to entities after delivery to the owner		X		X	X	
Update regularly the credentials to access its domain			X	X	X	X
Provide components or devices to owner without remote backdoor access	X	X				
Create / Modify / Delete a domain before delivery to owner		X				
Create / Modify / Delete a domain after delivery to owner				X	X	
Configure the access policy of the device ressources				X	X	
Enforce the device ressource access policy				X	X	
Configure its domain ressource access policy			X	X	X	X
Enforce its domain ressource access policy			X	X	X	X
Provide updates & patches of components	X					
Follow up, deploy updates of firmware or drivers			X			
Keep the service applications up to date						X

reducing the security of the system. Any entity which owns a partition (administrator, service provider, and maintainer) is responsible for modifying the temporary credentials of its associated domains.

C. Architecture requirements

1) Domain isolation: Two services which run in two isolated domains should not be able to access the memory of one another, gain information about each other or interfere with each other's execution.

2) Owner non-interference: Although the owner or the administrator have access to a privileged domain, they must not have any control or visibility on the actions performed by other entities unless these actions involve device ressources under the responsibility of the owner.

3) Ressource access control: A device ressource access policy under the control of the owner or the administrator is mandatory. For each domain, a domain ressource access policy under the control of the domain owner (maintainer, service provider) is optional.

4) Minimal on-device processing: The processing required to authorize an action in the device must be minimal.

V. PROPOSED SECURITY MODEL AND ARCHITECTURE

This section describes the architecture proposed in the ODSI project. First, we describe the overall architecture. Then, we define the notion of domain in this architecture as well as its components. Finally, we describe the communication and management models.

A. Definition of a domain

In this architecture, the device is shared between several parties with different roles (described in section IV-A). As one party can execute various pieces of code in different spaces, the following vocabulary is used:

- A partition is a piece of code managed by one party with its own memory space
- A domain is the union of all partitions managed and developped by one party.

Figure 5 illustrates an example of memory layout with 3 domains. The owner domain is the most privileged and the parent of domains 2 (maintainer) and 3 (service provider). In each domain, subpartitions allow to further isolate components. For example, in the owner domain (Dom0), the drivers (sensor driver, actuator driver, network driver) and the managers which allow to virtualise them are isolated from each others. They are also isolated from the scheduler, the configuration manager and the components Admin Manager, Key Vault, Token & Security Validator and Internal Communicator which enforce the security management model described in this article.

B. Components of a domain

This section describes the role of the components which enforce the proposed security model, namely the Configuration Manager, the Administration Manager, the Token & Security Validator and the Key Vault. The architecture proposed is also based on the use of virtual sensors and actuators.

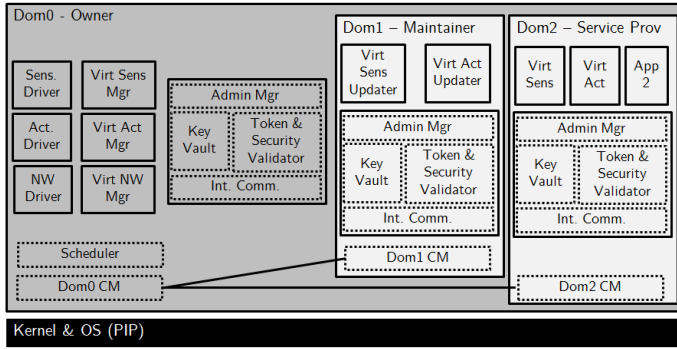Figure 6 represents the software layers which correspond to the example used in figure 5. The PIP kernel exposes

Fig. 5. Proposed architecture - Memory layer view

isolation services to an OS on which domains and partitions are created. The OS layer contains a scheduler which organizes the time and operations of the domains and partitions and each domain configuration manager. The application layer contains the drivers and all the modules of the architecture defined here. The OS and application layers both belong to user land.
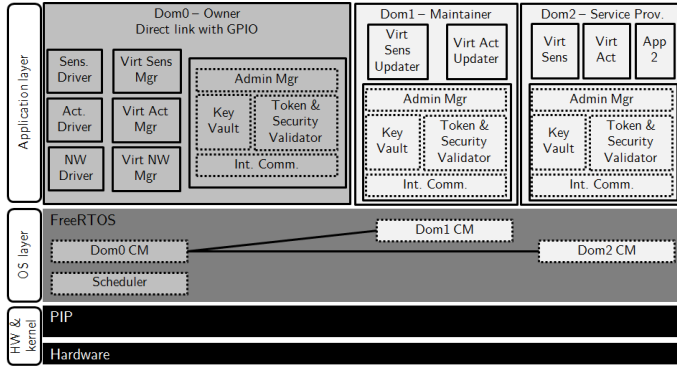

Fig. 6. Proposed architecture - Software layer view

1) Configuration Manager: The Configuration Manager contains a description of the domain's sub partitions, the associated code and the authorized communication channels between two partitions. It enforces the communications mechanism based on shared memory and described in section V-D.

2) Internal Communicator: The Internal Communicator is the unique entry point of the domain. It verifies incoming messages whether they come from the network or another domain, similarly to a firewall.

3) Virtual sensors or actuators managers: These managers provides synthetic sensor and actuator support to the virtual sensors and actuators over the internal communication module.

4) Virtual sensors or actuators: Only the owner domain can manage the machine-level functions such as the drivers or the physical memory. Each domain contains a virtual sensor or actuator which are a synthetic sensor or actuator instance. It consumes the hardware ressources provided by the virtual sensor or actuator manager and requires and forwards a token to the owner domain if required. This allows the entity which owns the domain to see his domain as an actual device without

knowledge of the isolation. The instruction and response are transmitted to and from the real driver using the internal communication mechanism. The virtual sensor or actuator updater is specific and allows the update of the related sensor or actuator driver in the owner domain.

5) Administration Manager: This module exposes the domain's ressources to an external server managed by the entity which owns the domain. It routes the command to read, write, execute the ressource to the expected manager or virtual sensor, actuator. It sends each command received to the Token & Security Validator (T&S Validator).

6) Token & Security Validator: The Token & Security Validator (T&S Validator) validates each device management command against the token provided in the command and optionnally an internal security policy. For example, the security policy defines which ressources require a token verification to be accessed or whether a specific command can be processed accordaccording to the device battery status.

7) Key Vault: This module stores the keys that are used by the domain. In particular, the keys used by the Token & Security Valdiator are stored here. It also provides the functions to add a new key, modify or delete an existing key.

C. Management model

As depicted in figure 7, the model proposed is based on master/agent model. Each domain receives commands from a server which authenticates the origin of the command and authorizes it. The administration manager in the domain is an agent which expects requests from its master, the administration application in the server. The Token & Security Validator validates the token generated according to the access policy on the server side. For simplicity's sake, only the components & modules related to the owner domain are detailed in figure 7.
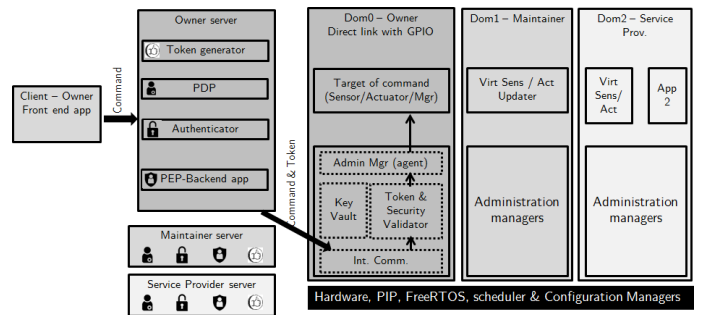

Fig. 7. Proposed architecture - Management model

Each entity is responsible for managing and authorizing the access of its own domain ressources (e.g. domain token keys, application installation) to its users. The owner also manages and authorizes the access to device ressources (e.g. sensor drivers & virtual managers) to its users and to other entities (e.g. maintainer or service provider).

1) Domain ressources management: An example is the modification of the URL from which to retrieve the package which contains a driver update. This parameter is in the perimeter of the maintainer and does not require any modification on the drivers or parameters controled by the owner. Therefore, it is the maintainer's responsiblity to ensure that this modification can be performed only by authenticated & authorized users.

As illustrated in figure 8, first, the client which requests the modification is authenticated. Then, it choses an operation proposed by the administration manager and the access policy evaluates whether this particular user is authorized to perform the selected operation. If yes, the token A is generated and the maintainer sends the command and the token are sent to the maintainer's domain where the command is processed.
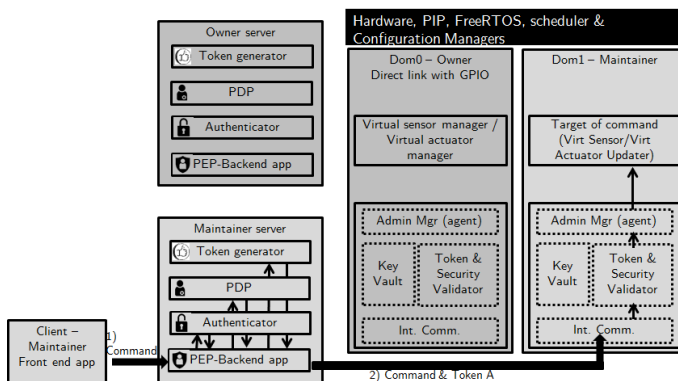

Fig. 8. Management of domain ressources

2) Device ressources management: An example is the download, deployment & installation of a package with a driver udpate. In this case, the maintaier authenticates the user which requests this operation and then requests the owner the authorization to perform it.

Figure 9 illustrates the steps necessary for this operation. First, the client which requests the update is authenticated and authorized by the maintainer server. Then, the maintainer server forwards the request to the owner. The owner authenticates the maintainer server which initiates the request and which will deploy the update package, authorizes the operation according to its own access policy and generates token B. The maintainer server now sends the command, token A and token B to the maintainer's domain. Once token A is verified by the maintainer's Token & Security Validator, the virtual sensor updator sends the command and token B to the Owner domain. If the owner validates token B, the update is authorized.

3) Token scheme: The token format should contain at minimum a payload and a signature of this payload. Compatible token formats include JWT (JSON Web Token) [6] and CWT (CBOR Web Token) [7].

The token's objectives are to 1) ensure that any command will reach the intended destination, 2) give the rights to
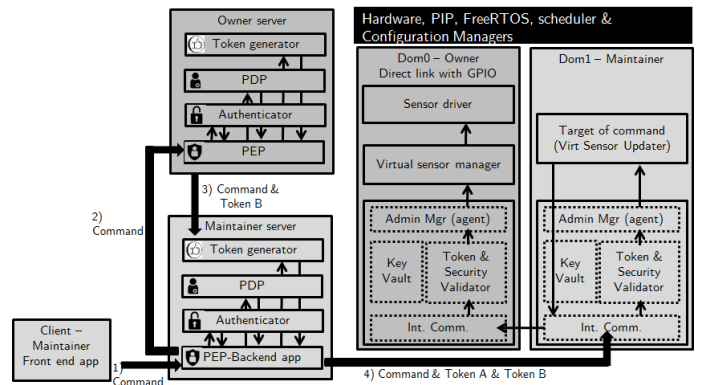

Fig. 9. Management of device ressources

request some operations on ressources, 3) prove that the entity initiating the command has been authenticated authorized and protect against replay. To achieve these goals, the token should contain at minimum:

- the ID of the partition to which this token is destined,
- the rights granted,
- a proof ensuring that the token was generated by the token provider,
- a token ID and/or an expiration date.

The proof can either be a digital signature, such as ANS X9.31, ANS X9.62 or PKCS#1 as recommended by NIST [8] or a Keyed-hash Message Authentication Codes (HMACS) compliant NIST recommendations and using a NIST-approved hash algorithm such as SHA-256 [9], [10], depending on the level of security and non-repudiation required. The token ID and expiration date are used to prevent replay. The token ID can be compared to a set of previously used IDs stored in a cache. The size of the cache should be adapted to the memory constraints of the partition and device. The combination of the size of the cache and the expiration date allows achieving best effort replay detection. Therefore, it is strongly recommended to use both elements and to tune them carefully.

D. Communication model

TBD. Section to be described by Lille

The main requirement for the communication protocol is to allow master/agent communication and the transport of a payload with the command content and an associated token. Protocols such as TCP/IP or CoAP are eligible. LwM2M [11] would require some modification to add a token to the message.

1) Internal communication:
2) External communication:

VI. PIP AND ODSI ARCHITECTURE

The aim of ODSI architecture is to provide a safe environment for each service provider. In other words, the aim of ODSI and Pip is how to preserve the trust between each service provider and between the platform and the service provider.

A. Trust between service providers

Each service provider has to trust the platform and the platform has to guarantee to each

> TOD0: - Confiance entre service, chaque service doit se voir garantir que son tat ne peut tre altr par un autre service, potentiellement concurent
> - Chaque service doit faire confiance la platforme, tant donne qu'elle est dans sa TCB.

B. Need in separation

> TOD0: - Pip fournit l'isolation recherche par les services pour garantir leur espace de travail/ espace d'adressage
> - Le mchanisme de token garantit que chaque service ne peut tre accd que via l'API qu'il aura expos et sans plus

VII. ILLUSTRATION & DISCUSSION

> TBD. Section to show how the isolation model allows to address the industrial needs

From the scenario described above [2], we can use our proposed ODSI security architecture and Pip to avoid this kind attack on any household devices.

We can take the example presented in figure 10. Each device is connected to a service provider hosted by the ODSI platform, and it has not a direct connexion to other devices. We can illustrate this architecture with figure 11. If the Smartphone wants to start a music on the music provider, it has to use the application available into the ODSI platform.
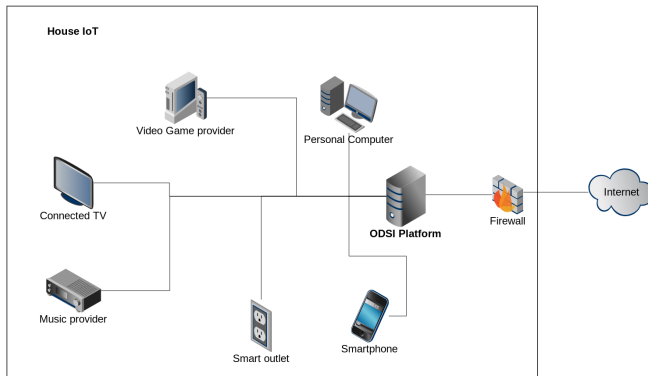

Fig. 10. House IoT installation example

Now, with the same attack scenario as described above, if the music application is not secure and can give access to attacker from outside the house, how the ODSI platform can prevent from any attack propagation ?

As explained, memory isolation provided by Pip has the ability to avoid for music app to access to other domain's memory spaces (Video game app, Owner domain...). The only way for this application to contact other available devices is to use the internal communication mechanism provided by Pip and the ODSI platform. The token mechanism allows only a set of authorized message to a specific service. This list of
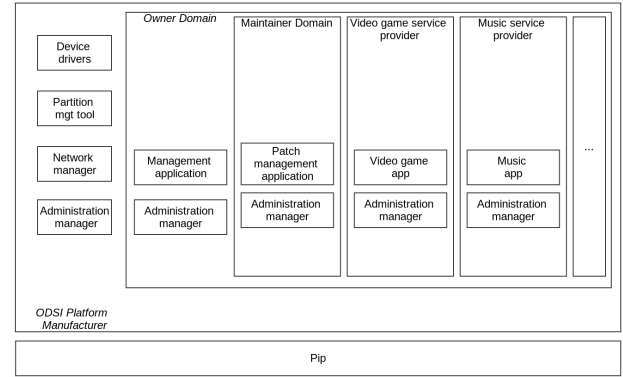

Fig. 11. ODSI Platform

these commands are provided by the service provider itself. For example, for the music service, it can only authorize three commands : find song/singer, play music, stop music. If any malicious application running into the ODSI platform, it can not send corrupted command to the music service. Moreover, any corrupted service can not access to the hardware due to Pip and it can not use the network devicein order to find more fallible devices.

VIII. CONCLUSION

> TBD. The conclusion goes here. Next steps : implementation and performance evaluation

REFERENCES

[1] Jo Ann Oravec. Kill switches, remote deletion, and intelligent agents: Framing everyday household cybersecurity in the internet of things. *Technology in Society*, 51:189–198, 2017.

[2] Vijay Sivaraman, Dominic Chan, Dylan Earl, and Roksana Boreli. Smart-Phones Attacking Smart-Homes. In *Proceedings of the 9th ACM Conference on Security & Privacy in Wireless and Mobile Networks - WiSec '16*, pages 195–200, New York, New York, USA, 2016. ACM Press.

[3] Quentin Bergougnoux, Julien Iguchi-Cartigny, and Gilles Grimaud. Pip, un proto-noyau fait pour renforcer la sécurité dans les objets connectés. In *Conférence d'informatique en Parallélisme, Architecture et Système (ComPAS)*, Sophia Antipolis, France, June 2017. Université Sophia Antipolis.

[4] ENISA. Baseline security recommendations for iot in the context of critical information infrastructures. Technical report, ENISA, 2017.

[5] Mark Warner, Gardner, Wyden, and Senate of the United States Daines. Internet of things (iot) cybersecurity improvement act of 2017. https://www.congress.gov/115/bills/s1691/BILLS-115s1691is.pdf, August 2017.

[6] M. Jones, J. Bradley, and N. Sakimura. Json web token (jwt). Technical report, IETF, 2015.

[7] M. Jones, E. Wahlstroem, S. Erdtman, and H. Tschofenigs. Cbor web token (cwt). draft version.

[8] Cameron F.Kerry and Patrick D.CGallagher. Digital signature standard. Technical report, National Institute of Standards and Technology, July 2013.

[9] Carlos Gutierrez and James M.Turner. The keyed-hash message authentication code (hmac). Technical report, National Institute of Standards and Technology, July 2008.

[10] Quynh Dang. Recommendation for applications using approved hash algorithms. Technical report, National Institute of Standards and Technology, 2012.

[11] OMA. Lightweight machine to machine technical specification. Technical report, Open Mobile Alliance, 2017.